

The MIPS Datapath

The Critical Path and Circuit Timing


- The critical path is the slowest path through the circuit
 - For a synchronous circuit, the clock cycle ***must be longer*** than the critical path otherwise we will have timing problems
- Its the combination of the CLK->Q time for the flip flop, the combinational delay, and the setup time
 - CLK->Q: Time it takes from when the clock edge comes to the flip-flop output changes
 - Combinational delay: The slowest/longest path in the combinational logic
 - Setup time: Time that the flip-flop requires the the input be stable before the clock comes

And Hold Time...

- Flip-flops are delicate little circuits...
 - They don't like their input changing before the clock comes (setup time)
 - And they don't like their input changing just after the clock edge (hold time)
- Hold time violations occur when there exists a path through the circuit that is very short
 - $\text{CLK} \rightarrow \text{Q} + \textit{shortest} \text{ combinational delay} < \text{hold time}$
 - Fix by deliberately adding in delays

How A D-type Flip-Flop Works

- A Flip-Flop is two transparent D-type latches
 - A latch passes $D \rightarrow Q$ when C is high
 - A latch holds Q when C is low




Latch Truth Table

D	CLK	Q
X	0	NC
0	1	0
1	1	1

This is why we need:
Setup time
Hold Time
CLK \rightarrow Q time

Arithmetic and Logic Unit

- Most processors contain a special logic block called the “Arithmetic and Logic Unit” (ALU)
- We’ll show you an easy one that does ADD, SUB, bitwise AND, bitwise OR




when $S=00$, $R=A+B$

when $S=01$, $R=A-B$

when $S=10$, $R=A \text{ AND } B$

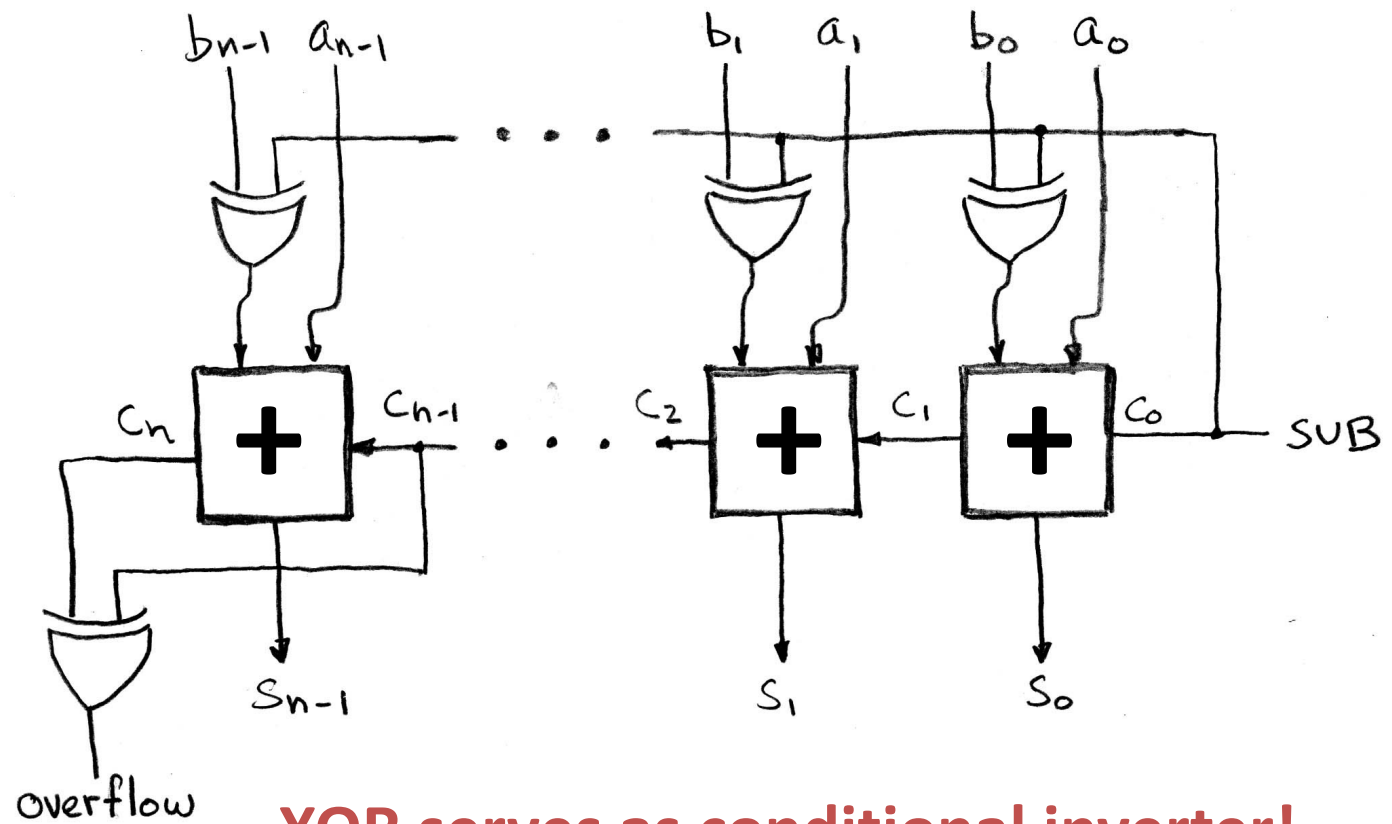
when $S=11$, $R=A \text{ OR } B$

Our simple ALU



Extremely Clever Subtractor

x	y	XOR(x,y)
0	0	0
0	1	1
1	0	1
1	1	0



XOR serves as conditional inverter!

Clickers – How was the midterm?

Emotion Spectrum



Clicker Question

Convert the truth table to a boolean expression
(no need to simplify):

A: $F = xy + x(\sim y)$

B: $F = xy + (\sim x)y + (\sim x)(\sim y)$

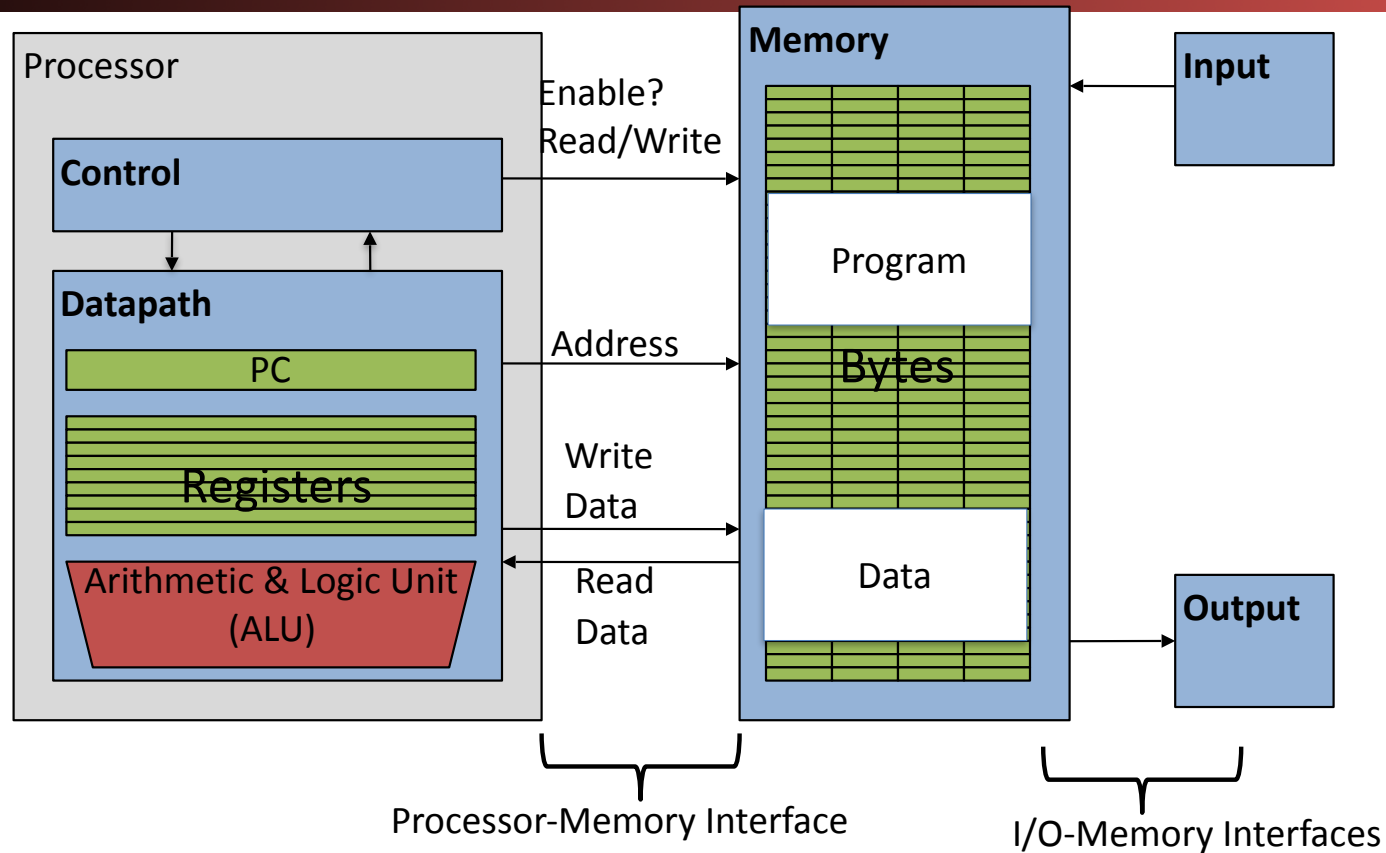
C: $F = (\sim x)y + x(\sim y)$

D: $F = xy + (\sim x)y$

E: $F = (x+y)(\sim x+\sim y)$

x	y	F(x,y)
0	0	0
0	1	1
1	0	0
1	1	1

Components of a Computer

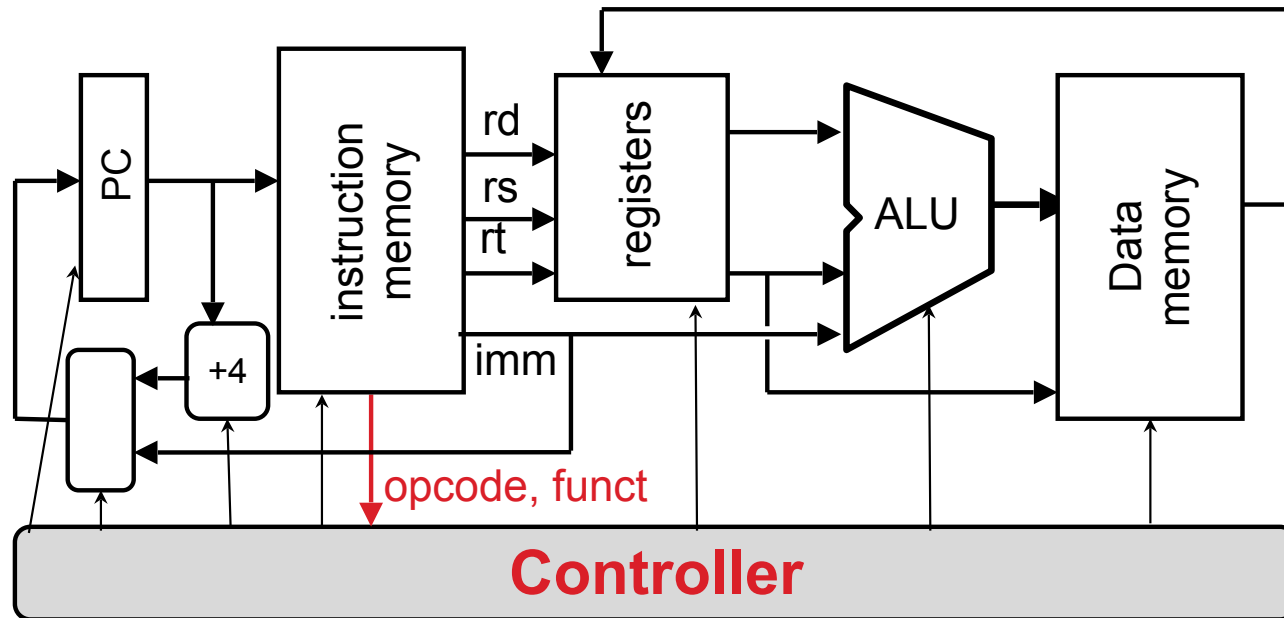


The CPU

- Processor (CPU): the active part of the computer that does all the work (data manipulation and decision-making)
- Datapath: portion of the processor that contains hardware necessary to perform operations required by the processor (the brawn)
- Control: portion of the processor (also in hardware) that tells the datapath what needs to be done (the brain)

Datapath and Control


- Datapath designed to support data transfers required by instructions
- Controller causes correct transfers to happen



Five Stages of Instruction Execution

- Stage 1: Instruction Fetch
- Stage 2: Instruction Decode
- Stage 3: ALU (Arithmetic-Logic Unit)
- Stage 4: Memory Access
- Stage 5: Register Write
- Later on we will add pipeline registers:
So when one instruction is in stage 3, the next instruction is in stage 2 and the one after that in stage 1

Stages of Execution on Datapath



Stages of Execution (1/5)

- There is a wide variety of MIPS instructions: so what general steps do they have in common?
- Stage 1: Instruction Fetch
 - no matter what the instruction, the 32-bit instruction word must first be fetched from memory (the cache-memory hierarchy)
 - also, this is where we Increment PC
(that is, $PC = PC + 4$, to point to the next instruction: byte addressing so + 4)
 - Which is why the branches are all relative to $PC + 4$, not PC

Stages of Execution (2/5)

- Stage 2: Instruction Decode
 - upon fetching the instruction, we next gather data from the fields (decode all necessary instruction data)
 - first, read the opcode to determine instruction type and field lengths
 - second, read in data from all necessary registers
 - for **add**, read two registers
 - for **addi**, read one register
 - for **jal**, no reads necessary

Stages of Execution (3/5)

- Stage 3: ALU (Arithmetic-Logic Unit)
 - the real work of most instructions is done here: arithmetic (+, -, *, /), shifting, logic (&, |), comparisons (slt)
 - what about loads and stores?
 - `lw $t0, 40($t1)`
 - the address we are accessing in memory = the value in \$t1 PLUS the value 40
 - so we do this addition in this stage

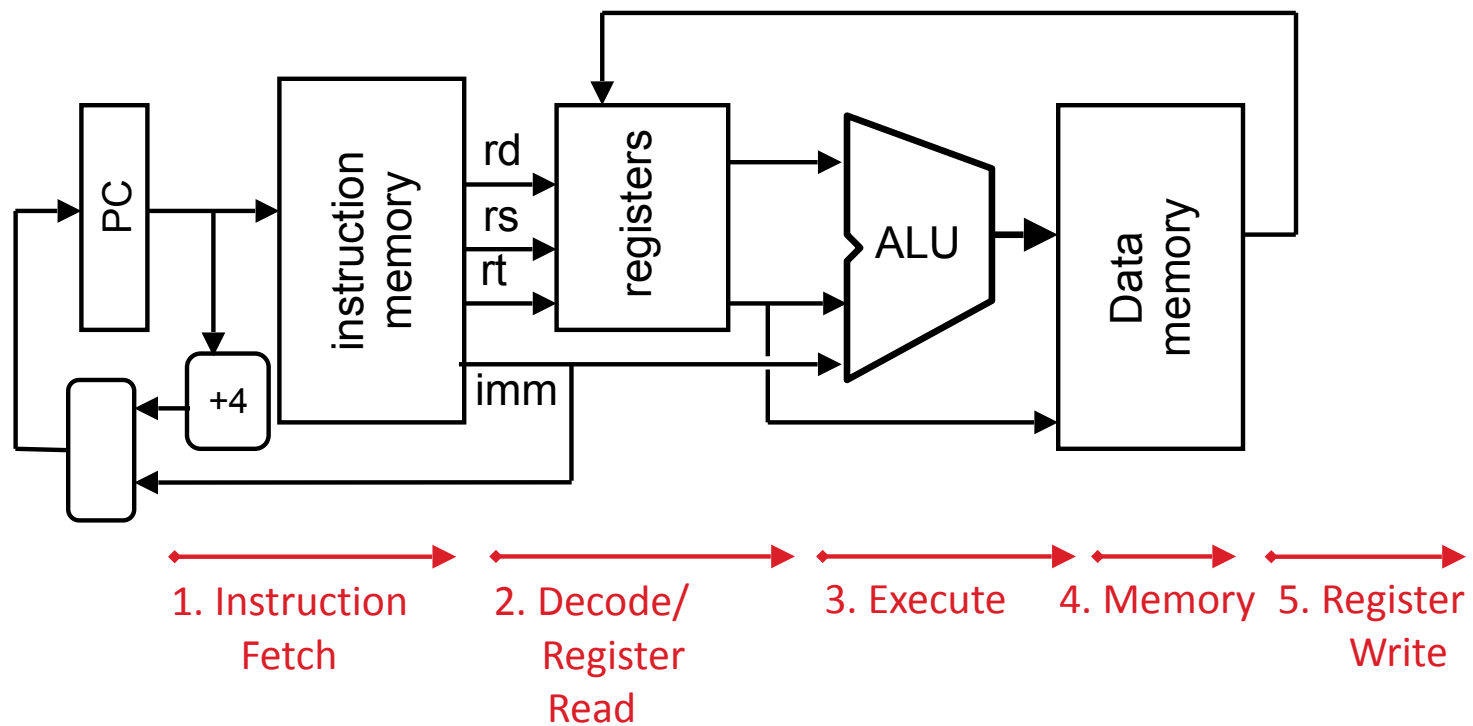
Stages of Execution (4/5)

- Stage 4: Memory Access
 - actually only the load and store instructions do anything during this stage; the others remain idle during this stage or skip it all together
 - since these instructions have a unique step, we need this extra stage to account for them
 - as a result of the cache system, this stage is expected to be fast
 - But if the cache fails to hit (more later in the class), this can "stall"

Stages of Execution (5/5)

- Stage 5: Register Write
 - most instructions write the result of some computation into a register
 - examples: arithmetic, logical, shifts, loads, slt
 - what about stores, branches, jumps?
 - don't write anything into a register at the end
 - these remain idle during this fifth stage or skip it all together


Stages of Execution on Datapath



Datapath Walkthroughs (1/3)

- **add \$r3,\$r1,\$r2 # r3 = r1+r2**
 - Stage 1: fetch this instruction, increment PC
 - Stage 2: decode to determine it is an add, then read registers **\$r1** and **\$r2**
 - Stage 3: add the two values retrieved in Stage 2
 - Stage 4: idle (nothing to write to memory)
 - Stage 5: write result of Stage 3 into register **\$r3**


Example: add Instruction



Datapath Walkthroughs (2/3)

- `slti $r3,$r1,17`
if (r1 < 17) r3 = 1 else r3 = 0
- Stage 1: fetch this instruction, increment PC
- Stage 2: decode to determine it is an `slti`, then read register `$r1`
 - We will also likely fetch `$r3` but value is ignored:
Simplifies decoding logic on deciding which registers to fetch
- Stage 3: compare value retrieved in Stage 2 with the integer 17
- Stage 4: idle
- Stage 5: write the result of Stage 3 (1 if reg source was less than signed immediate, 0 otherwise) into register `$r3`

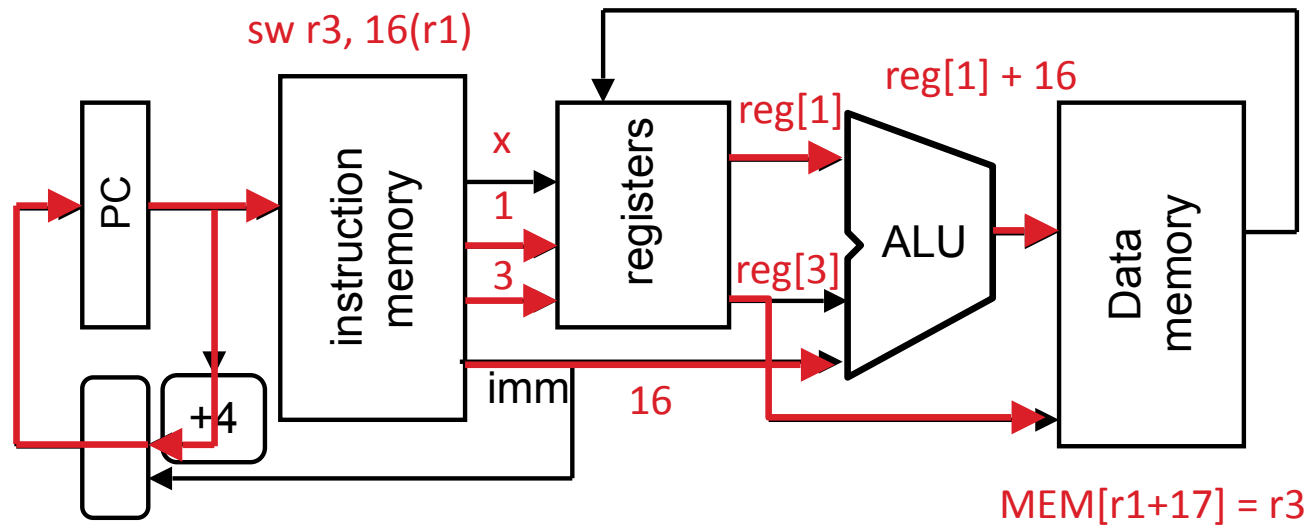
Example: slti Instruction



Datapath Walkthroughs (3/3)

- **sw \$r3,16(\$r1) # Mem[r1+16]=r3**
 - Stage 1: fetch this instruction, increment PC
 - Stage 2: decode to determine it is a sw, then read registers **\$r1** and **\$r3**
 - Stage 3: add 16 to value in register \$r1 (retrieved in Stage 2) to compute address
 - Stage 4: write value in register \$r3 (retrieved in Stage 2) into memory address computed in Stage 3
 - Stage 5: idle (nothing to write into a register)

Example: sw Instruction




Why Five Stages? (1/2)

- Could we have a different number of stages?
 - Yes, other ISAs have different natural number of stages
 - And these days, pipelining can be much more aggressive than the "natural" 5 stages MIPS uses
- Why does MIPS have five if instructions tend to idle for at least one stage?
 - Five stages are the union of all the operations needed by all the instructions.
 - One instruction uses all five stages: the load

Why Five Stages? (2/2)

- `lw $r3, 16($r1) # r3=Mem[r1+16]`
 - Stage 1: fetch this instruction, increment PC
 - Stage 2: decode to determine it is a `lw`, then read register `$r1`
 - Stage 3: add 16 to value in register `$r1` (retrieved in Stage 2)
 - Stage 4: read value from memory address computed in Stage 3
 - Stage 5: write value read in Stage 4 into register `$r3`

Example: lw Instruction



Processor Design: 5 steps

Step 1: Analyze instruction set to determine datapath requirements

- Meaning of each instruction is given by register transfers
- Datapath must include storage element for ISA registers
- Datapath must support each register transfer

Step 2: Select set of datapath components & establish clock methodology

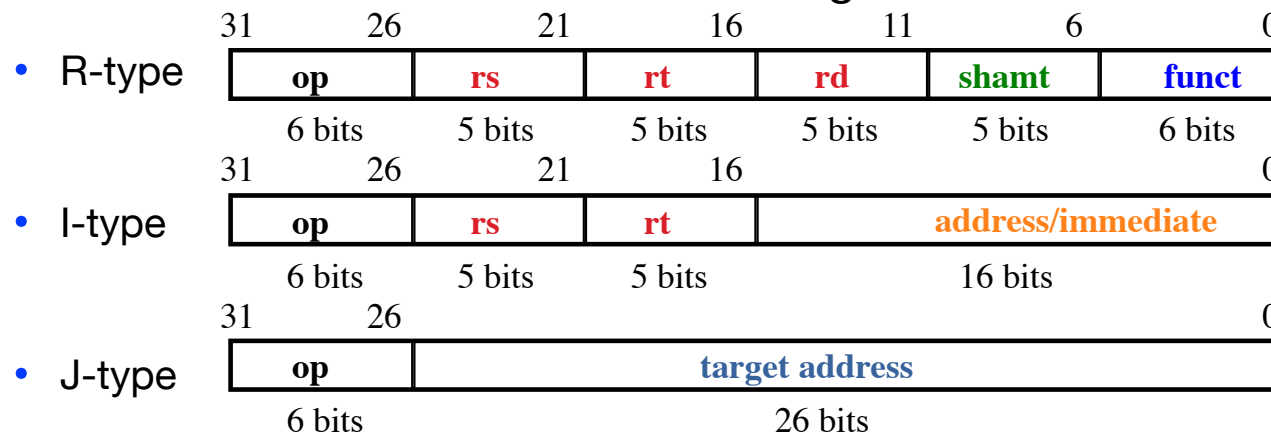
Step 3: Assemble datapath components that meet the requirements

Step 4: Analyze implementation of each instruction to determine setting of control points that realizes the register transfer

Step 5: Assemble the control logic

The MIPS Instruction Formats





- All MIPS instructions are 32 bits long. 3 formats:



- The different fields are:
 - op**: operation (“opcode”) of the instruction
 - rs, rt, rd**: the source and destination register specifiers
 - shamt**: shift amount
 - funct**: selects the variant of the operation in the “op” field
 - address / immediate**: address offset or immediate value
 - target address**: target address of jump instruction

The MIPS-lite Subset

- ADDU and SUBU
 - `addu rd,rs,rt`
 - `subu rd,rs,rt`
- OR Immediate:
 - `ori rt,rs,imm16`
- LOAD and STORE Word
 - `lw rt,rs,imm16`
 - `sw rt,rs,imm16`
- BRANCH:
 - `beq rs,rt,imm16`



Register Transfer Level (RTL)

- Colloquially called “Register Transfer Language”
- RTL gives the meaning of the instructions
- All start by fetching the instruction itself

$\{op, rs, rt, rd, shamt, funct\} \leftarrow MEM[PC]$

$\{op, rs, rt, Imm16\} \leftarrow MEM[PC]$

Inst Register Transfers

ADDU $R[rd] \leftarrow R[rs] + R[rt]; PC \leftarrow PC + 4$

SUBU $R[rd] \leftarrow R[rs] - R[rt]; PC \leftarrow PC + 4$

ORI $R[rt] \leftarrow R[rs] \mid \text{zero_ext}(Imm16); PC \leftarrow PC + 4$

LOAD $R[rt] \leftarrow MEM[R[rs] + \text{sign_ext}(Imm16)]; PC \leftarrow PC + 4$

STORE $MEM[R[rs] + \text{sign_ext}(Imm16)] \leftarrow R[rt]; PC \leftarrow PC + 4$

BEQ $\text{if } (R[rs] == R[rt])$
 $PC \leftarrow PC + 4 + \{\text{sign_ext}(Imm16), 2'b00\}$
 else $PC \leftarrow PC + 4$

In Conclusion

- “Divide and Conquer” to build complex logic blocks from smaller simpler pieces (adder)
- Five stages of MIPS instruction execution
- Mapping instructions to datapath components
- Single long clock cycle per instruction

If There is Time Left: CloudBleed...

- How many got "change your password" emails from various services?
- What happened: C out-of-bounds read!
 - `if(ptr == end) ...`
 - Unfortunately, `ptr` got incremented 1 past end when it couldn't parse HTML right
 - Result was that Cloudflare would, under some conditions, read out memory
- This memory contained other requests to ***different*** cloudflare sites
 - And it would get picked up by search engines and stored in their caches!

Lessons Learned...

- Without detailed audibility, its impossible to know what was lost
- Which is why some sites went "oh, change your password", even though the odds of your password being compromised are slim

Lessons Already Forgotten...

- Don't Use C/C++!
 - This is "yet another massively disruptive bug caused because programmers don't use safe languages!"
- And if you do, don't use pointer arithmetic
 - Array logic, eg `if (index >= end) ...` rather less likely to screw up
- Use Valgrind or similar
 - Valgrind would have trapped this as an "out of bounds read"